



# An Introduction to Deep Reinforcement Learning

Ehsan Abbasnejad



# Remember: Supervised Learning

We have a set of sample observations, with **labels**

learn to predict the labels, given a new sample



cat



dog

Learn the function that  
associates a picture of a  
dog/cat with the label

# Remember: supervised learning

We need thousands of samples

Samples have to be provided by experts

There are applications where

- We can't provide expert samples
- Expert examples are not what we mimic
- There is an interaction with the world

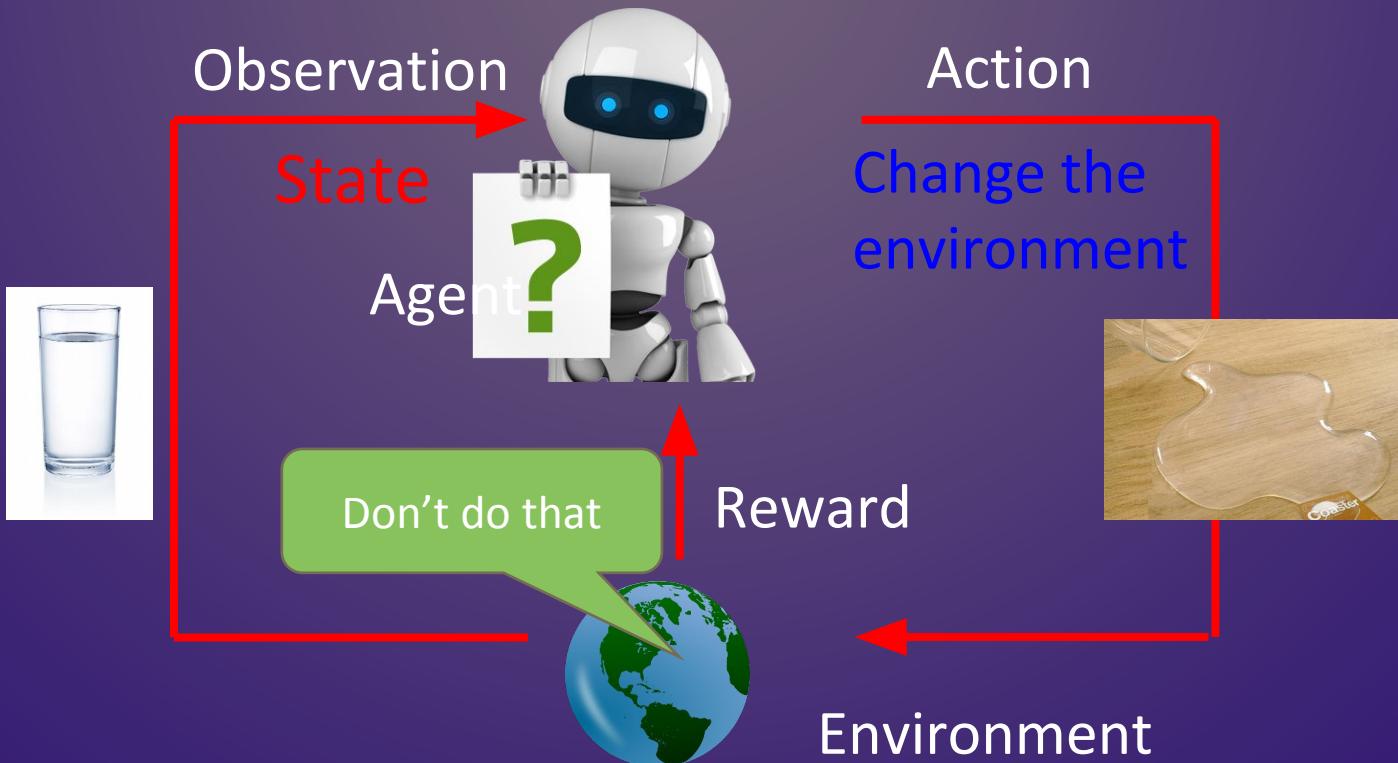
# Deep Reinforcement Learning



# AlphaGo



# Scenario of Reinforcement Learning



# Scenario of Reinforcement Learning

Agent learns to take actions maximizing expected reward.

Observation

State

Agent?

Action

Change the environment



Thank you.

Reward



Environment



# Machine Learning ≈ Looking for a Function

Actor/Policy

Action =  $\pi(\text{Observation})$

Observation

Function  
input



Action

Function  
output

Used to pick the  
best function

Reward



Environment

# Reinforcement Learning in a nutshell

RL is a general-purpose framework for decision-making

- RL is for an **agent** with the capacity to **act**
- Each **action** influences the agent's future **state**
- Success is measured by a scalar **reward** signal

Goal: **select actions to maximise future reward**

# Deep Learning in a nutshell

DL is a general-purpose framework for representation learning

- Given an **objective**
- Learning **representation** that is required to achieve objective
- Directly from **raw inputs**
- Using minimal domain knowledge

Goal: Learn the representation that achieves the  
objective

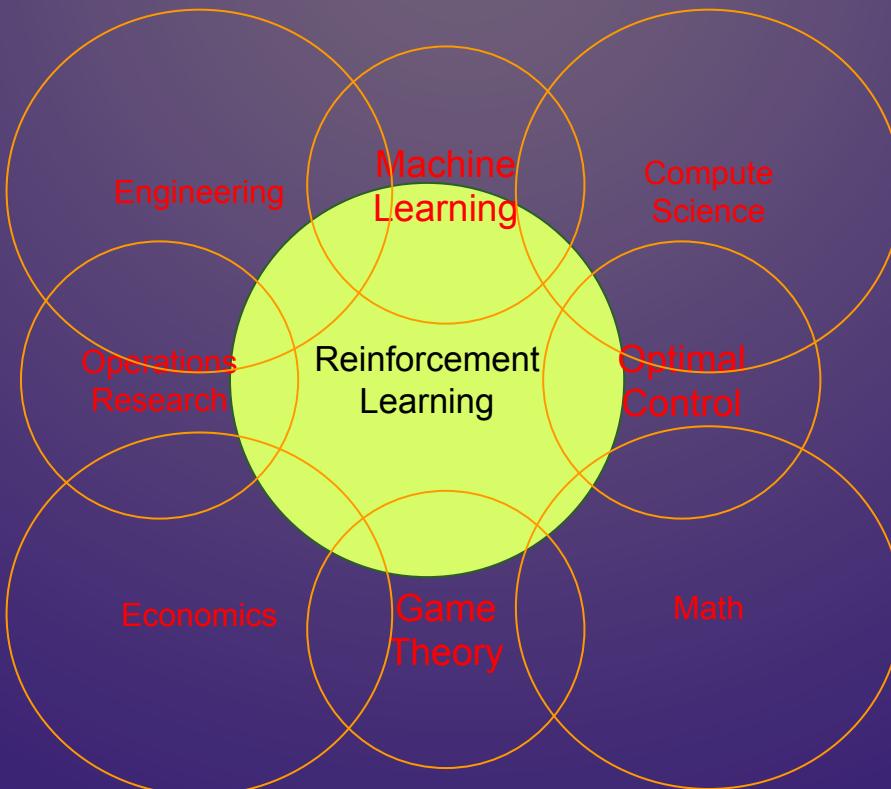
# Deep Reinforcement Learning in a nutshell

A single agent that solves human level tasks

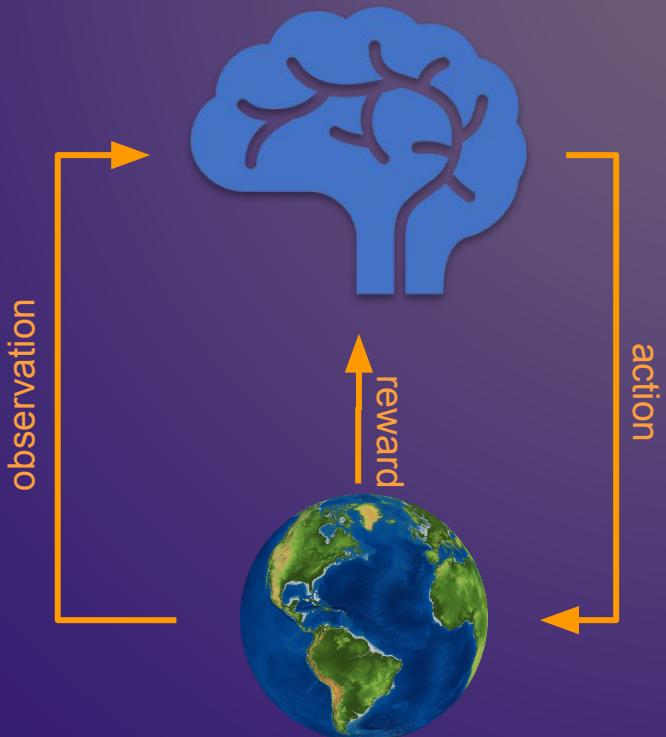
- RL defines the objective
- DL gives the mechanism and representation
- RL+DL=Deep reinforcement learning

This can lead to general intelligence

# Reinforcement Learning is multi-disciplinary

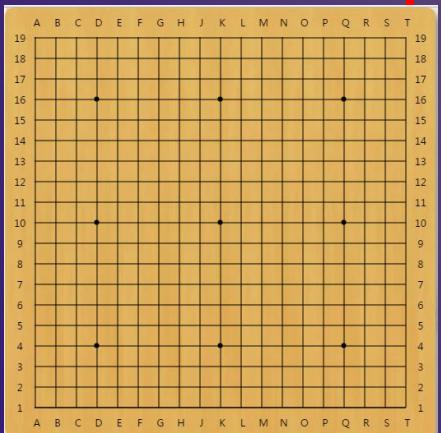


# Agent and Environment



- At each step, the agent
  - Selects an action
  - Observes the environment
  - Receives reward
- The environment:
  - Receives action
  - Emits new observation
  - Emits reward for the agent

# Learning to play Go

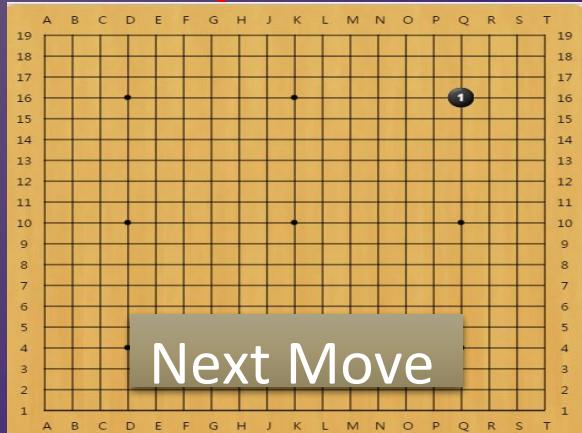


Observation



Reward

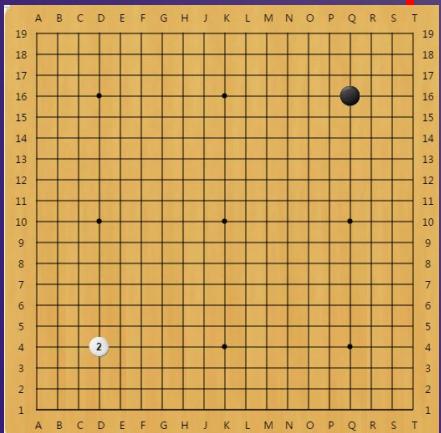
Action



Environment

# Learning to play Go

Agent learns to take actions maximizing expected reward.



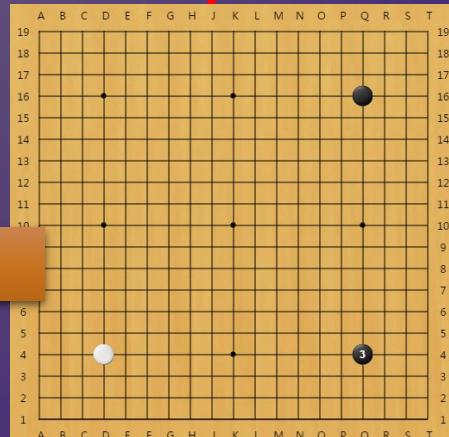
Observation



reward = 0 in most cases

If win, reward = 1

If loss, reward = -1



Action

Environment

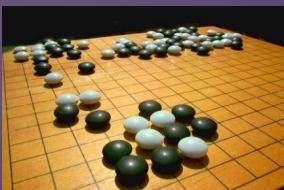
# Learning to play Go

- Supervised:

Learning from teacher



Next move:  
“5-5”



Next move:  
“3-3”

- Reinforcement Learning

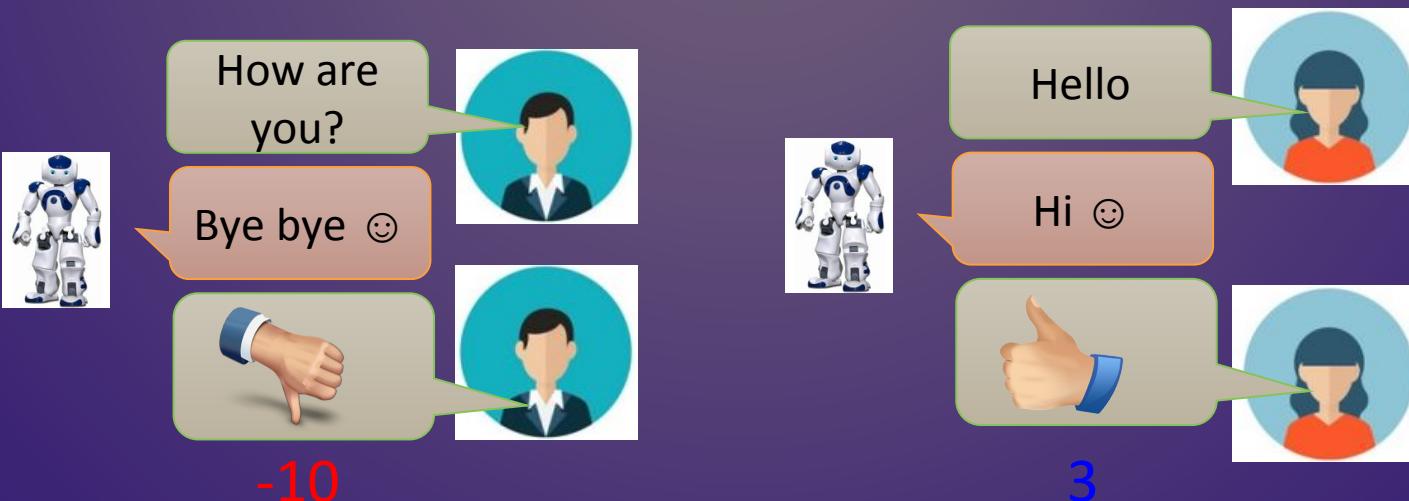
Learning from experience

First move → ..... many moves → Win!  
(Two agents play with each other.)

Alpha Go is supervised learning + reinforcement learning.

# Learning a chat-bot

- Machine obtains feedback from user



- Chat-bot learns to maximize the expected reward

[https://image.freepik.com/free-vector/variety-of-human-avatars\\_23-2147506285.jpg](https://image.freepik.com/free-vector/variety-of-human-avatars_23-2147506285.jpg)  
[http://www.freepik.com/free-vector/variety-of-human-avatars\\_766615.htm](http://www.freepik.com/free-vector/variety-of-human-avatars_766615.htm)

# Learning a chat-bot

- Let two agents talk to each other (sometimes generate good dialogue, sometimes bad)



How old are you?



How old are you?



I am 16.



See you.



See you.



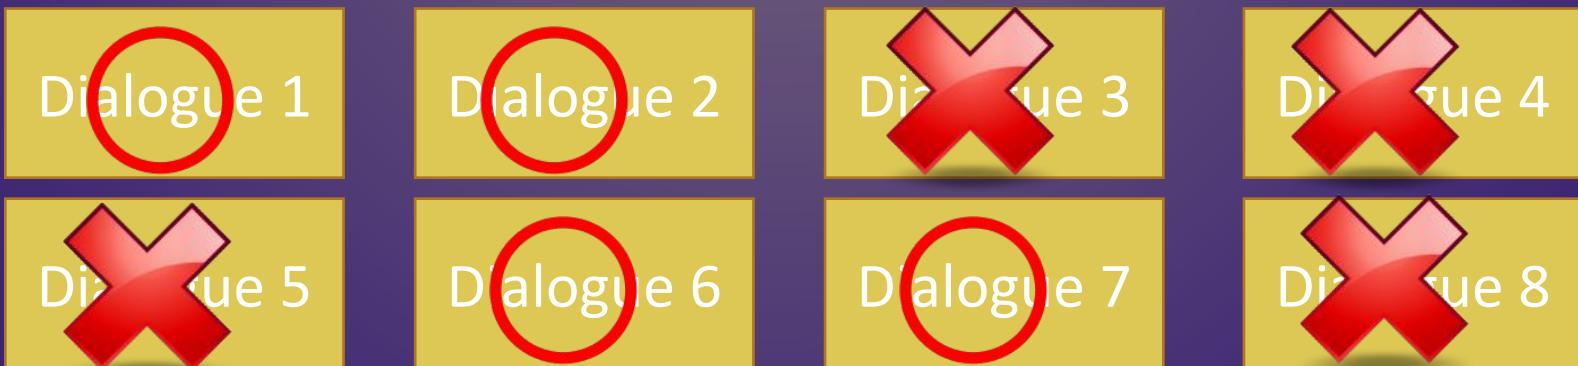
I thought you were 12.



What make you  
think so?

# Learning a chat-bot

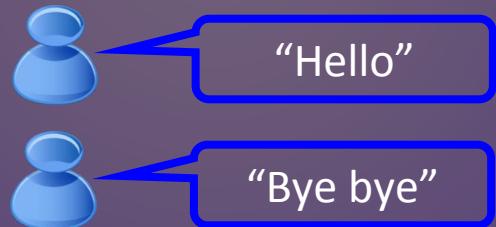
- By this approach, we can generate a lot of dialogues.
- Use some predefined rules to evaluate the goodness of a dialogue



Machine learns from the evaluation

# Learning a chat-bot

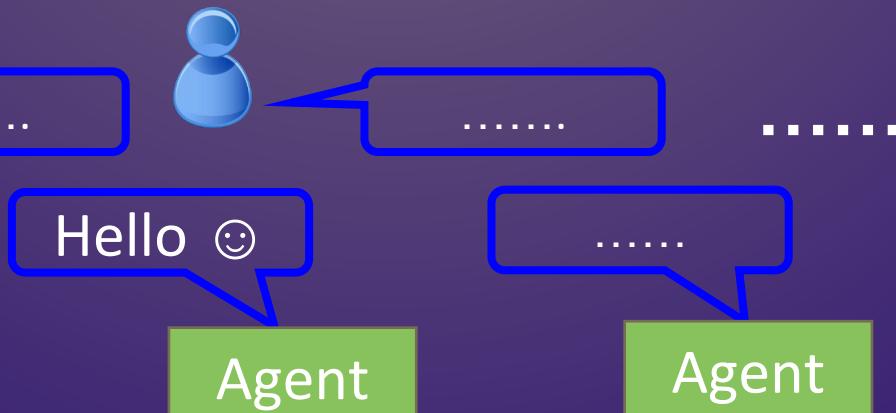
- Supervised



Say "Hi"

Say "Good bye"

- Reinforcement



Bad

# More applications

- Flying Helicopter
  - <https://www.youtube.com/watch?v=0JL04JJjocc>
- Driving
  - <https://www.youtube.com/watch?v=0xo1Ldx3L5Q>
- Robot
  - <https://www.youtube.com/watch?v=370cT-OAzzM>
- Google Cuts Its Giant Electricity Bill With DeepMind-Powered AI
  - <http://www.bloomberg.com/news/articles/2016-07-19/google-cuts-its-giant-electricity-bill-with-deepmind-powered-ai>
- Text generation
  - <https://www.youtube.com/watch?v=pbQ4qe8EwLo>

# Example: Playing Video Game

- Widely studies:
  - Gym: <https://gym.openai.com/>
  - Universe: <https://openai.com/blog/universe/>

Machine learns to play video  
games as human players

- What machine observes is pixels
- Machine learns to take proper action itself

## Example: Playing Video Game

Termination: all the aliens are killed, or your spaceship is

- Space invader



## Example: Playing Video Game

- Space invader
- Play yourself:  
<http://www.2600online.com/spaceinvaders.html>
- How about machine:  
[https://gym.openai.com/evaluations/eval\\_Eduozx4HRyqgTCVk9ltw](https://gym.openai.com/evaluations/eval_Eduozx4HRyqgTCVk9ltw)

# Example: Playing Video Game

Start with  
observation  $s_1$



Observation  $s_2$



Observation  $s_3$



Obtain reward  
 $r_1 = 0$

Action  $a_1$ : "right"



Obtain reward  
 $r_2 = 5$

Action  $a_2$  : "fire"  
(kill an alien)

Usually there is some randomness in the environment

# Example: Playing Video Game

Start with  
observation  $s_1$



Observation  $s_2$



Observation  $s_3$



After many turns



Obtain reward  $r_T$

Game Over  
(spaceship destroyed)

Action  $a_T$

This is an episode.

Learn to maximize the  
expected cumulative reward  
per episode

## Paradigm



Supervised  
Learning



Unsupervised  
Learning



Reinforcement  
Learning

## Objective

$$p_{\theta}(y|x)$$

$$p_{\theta}(x)$$

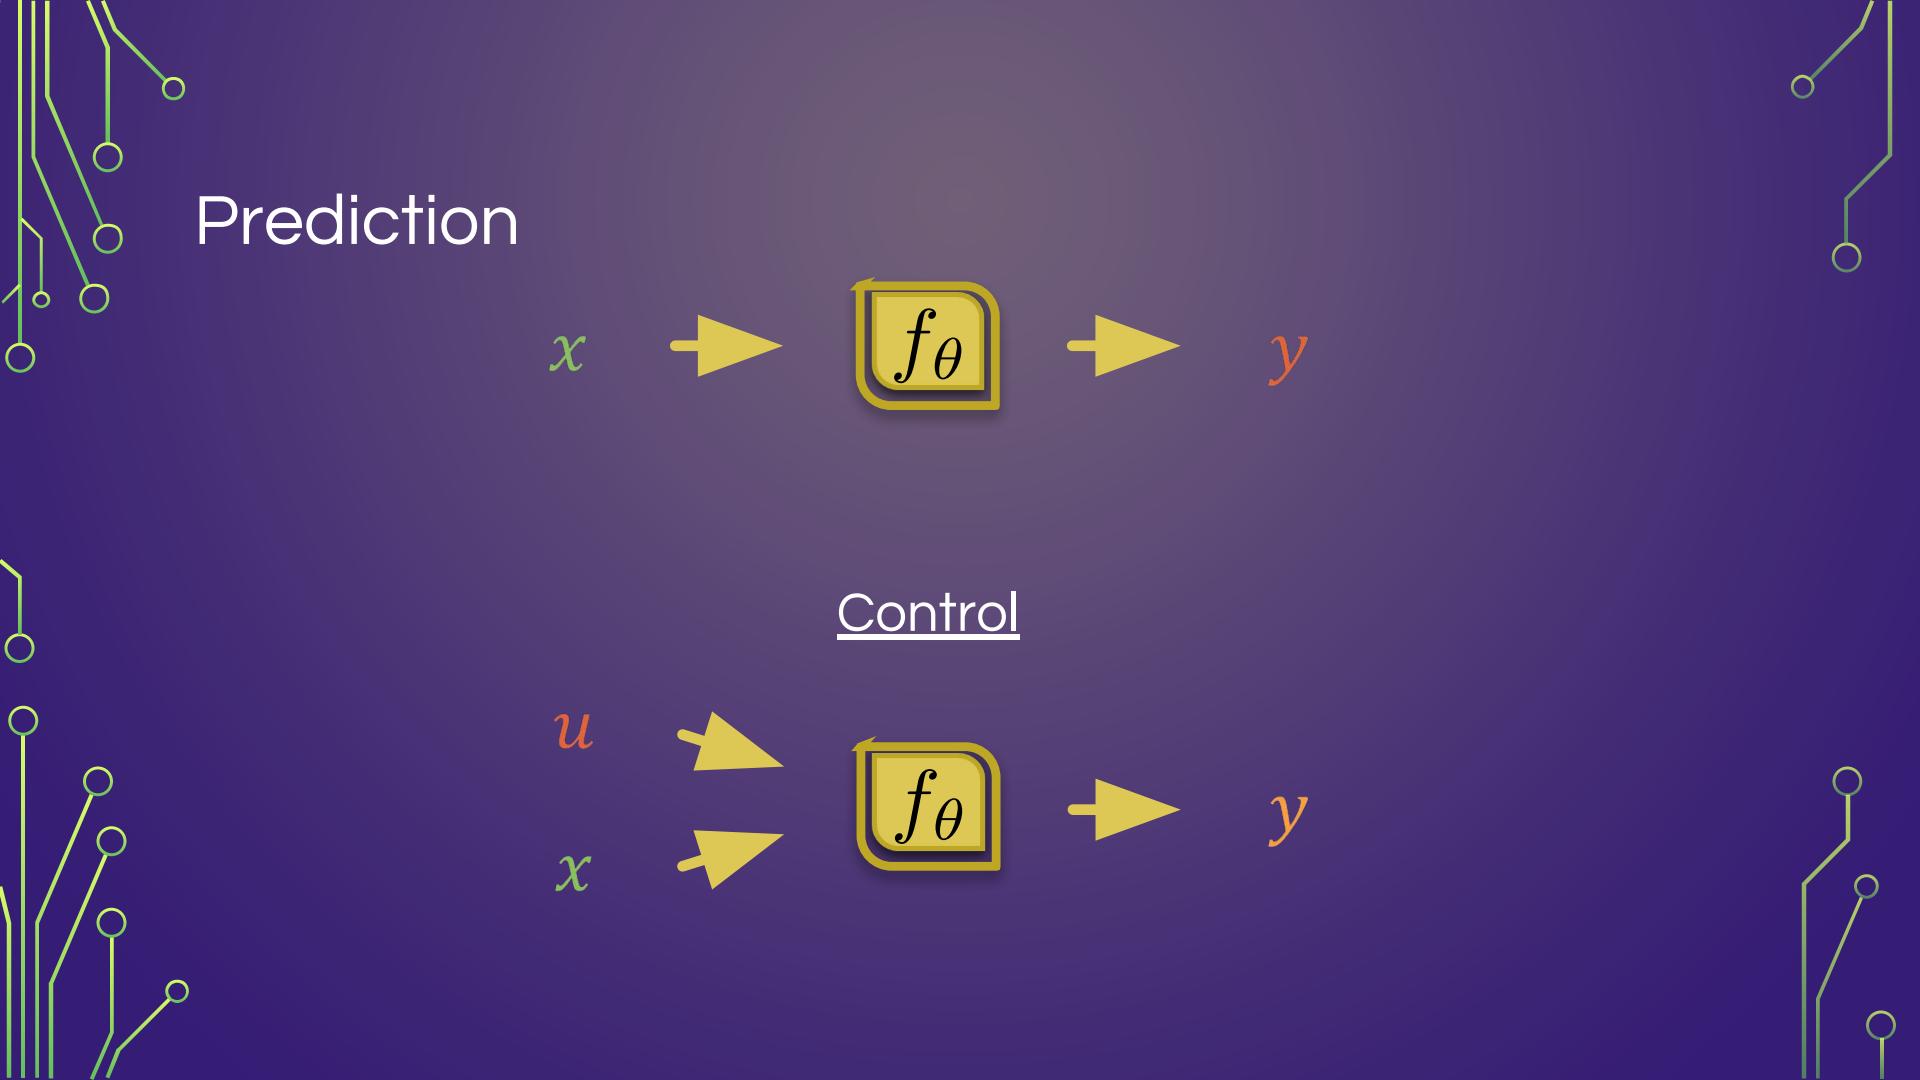
$$\pi_{\theta}(a|s)$$

## Applications

→ Classification  
→ Regression

→ Inference  
→ Generation

→ Prediction  
→ Control



Prediction



Control



SETTING

Environment

State/Observation  
Reward

Action

Agent

using  
policy

$\pi_\theta(a|s)$

# MARKOV DECISION PROCESSES (MDP)



State  
space



Action  
space



Transition  
function



Reward  
function

- **State:** Markov property considers only the previous state
- **Decision:** agent takes actions, and those decisions have consequences
- **Process:** there is a transition function (dynamics of the system)
- **Reward:** depends on the state and action, often related to the state

Goal: maximise overall reward

# Partially Observable MARKOV DECISION PROCESSES (POMDP)



State space



Action space



Transition function



Reward function

- State: Markov property considers only the previous state **but the agent cannot directly observe the underlying state.**
- Decision: agent takes actions, and those decisions have consequences
- Process: there is a transition function (dynamics of the system)
- Reward: depends on the state and action, often related to the state

Goal: maximise overall reward

# MARKOV DECISION PROCESSES (MDP)



State  
space



Action  
space



Transition  
function



Reward  
function

$$s_t \in \mathcal{S}$$

$$a_t \in \mathcal{A}$$

$$\mathcal{T} : \mathcal{S} \times \mathcal{A} \mapsto \mathcal{S}$$

$$s_{t+1} \sim \mathcal{T}(\cdot | s_t, a_t)$$

$$s_0 \sim \mathcal{T}_0$$

$$\mathcal{R} : \mathcal{S} \times \mathcal{A} \mapsto \mathbb{R}$$

$$r_t \sim \mathcal{R}(s_t, a_t)$$

# Computing Rewards

Episodic vs continuing: “Game over” after N steps

Additive rewards (can be infinite for continuing tasks)

Discounted rewards ...

# DISCOUNT FACTOR

- We want to be **greedy** but not **impulsive**
- Implicitly takes uncertainty in dynamics into account (we don't know the future)
- Mathematically:  $\gamma < 1$  allows infinite horizon returns

Return: 
$$G(s_t, a_t) = \sum_{\tau=0}^T \gamma^\tau \mathcal{R}(s_{t+\tau}, a_{t+\tau})$$

# SOLVING AN MDP

Objective:

$$J(\pi) = \mathbb{E}_{a_t \sim \pi(\cdot | s_t), s_{t+1} \sim \mathcal{T}(\cdot | s_t, a_t), s_0 \sim \mathcal{T}_0} \left[ \sum_{t=0}^T \gamma^t \mathcal{R}(s_t, a_t) \right]$$

Goal:  $\hat{\pi} = \arg \max_{\pi} J(\pi)$

# SOLVING AN MDP

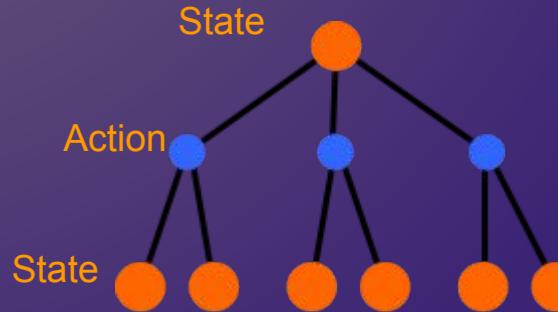
- If the state and actions are discrete:
  - We have a table of state-action probabilities
  - Learning is filling this table: (dynamic programming)

Action			
State			

# SOLVING AN MDP

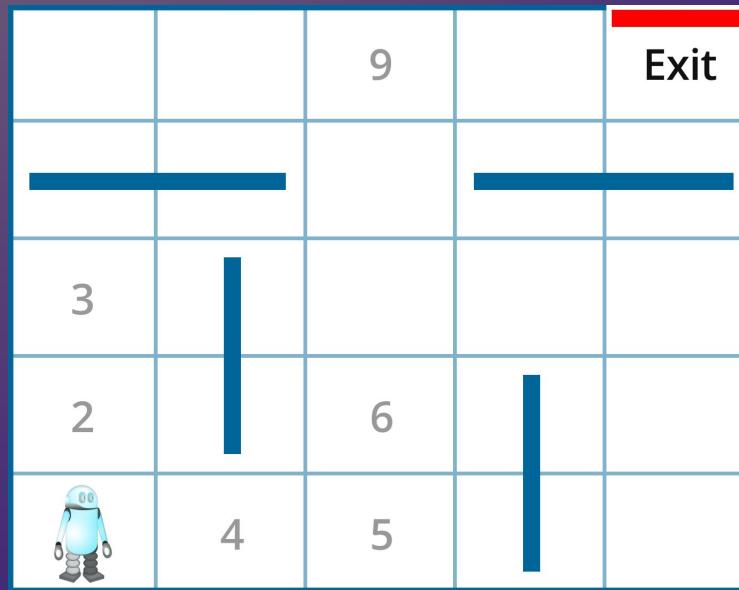
- If the state and actions are discrete:
  - We have a table of state-action probabilities
  - Learning is filling this table: (dynamic programming)

Action			



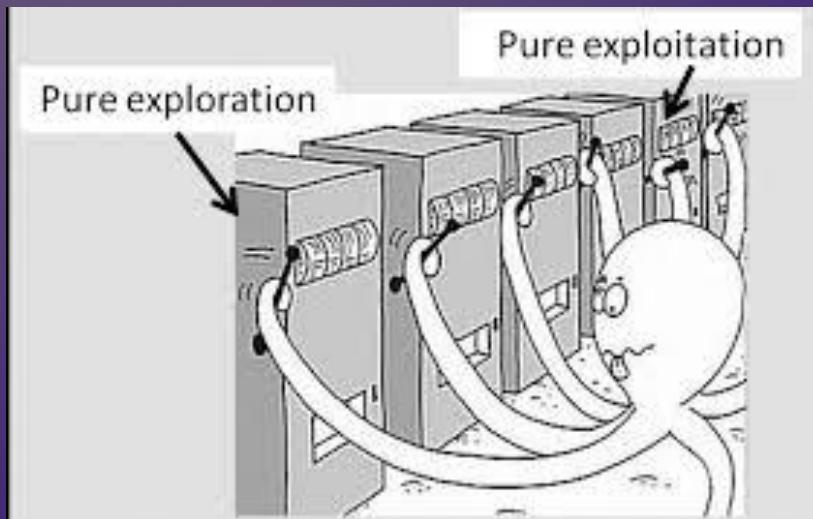
# SOLVING AN MDP

- If the state and actions are discrete:
- Let's try different actions and see which one succeed



# Exploration-Exploitation dilemma

Do we want to stick to action we think  
would be good or try something new



# Choosing Actions

- Take the action with highest probability (Q-function): Greedy
- Proportionate by its probability: Sampling
- Greedy most times, with some probability random

## VALUE FUNCTIONS

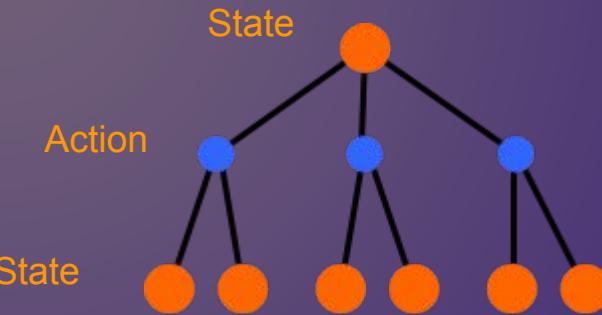
- Value = expected gain of a state
- Q function – action specific value function
- Advantage function – how much **more** valuable is an action
- Value depends on future rewards → depends on **policy**

$$V^\pi(s) = \mathbb{E}_{a_t \sim \pi(\cdot | s_t)} [G(s_0, a_0) | s_0 = s]$$

$$Q^\pi(s, a) = \mathbb{E}_{a_t \sim \pi(\cdot | s_t)} [G(s_0, a_0) | s_0 = s, a_0 = a]$$

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$$

## VALUE FUNCTIONS



$$V^\pi(s) = \mathbb{E}_{a_t \sim \pi(\cdot | s_t)} [G(s_0, a_0) | s_0 = s]$$

$$Q^\pi(s, a) = \mathbb{E}_{a_t \sim \pi(\cdot | s_t)} [G(s_0, a_0) | s_0 = s, a_0 = a]$$

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$$

# Solving Reinforcement Learning

- Model-based approaches:
  - We model the environment. Do we really need to model all the details of the world?
- Model free approaches:
  - We model the state-actions

Alpha Go: policy-based + value-based + model-based

## Model-free Approach

Policy-based

Value-based

Learning an Actor

Actor + Critic

Learning a Critic

## Model-based Approach

# POLICY ITERATION



Policy  
Evaluation



Policy  
Update

$$V^\pi(s) \leftarrow \sum_{a \in \mathcal{A}} \pi(a|s) Q^\pi(s, a)$$

$$Q^\pi(s, a) \leftarrow \pi(a|s) \left( \mathcal{R}(s, a) + \gamma \sum_{s' \in \mathcal{S}} \mathcal{T}(s'|s, a) V^\pi(s') \right)$$

$$\pi(s) \leftarrow \arg \max_a Q^\pi(s, a)$$

# Q-LEARNING

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha (r_t + \gamma Q(s_{t+1}, \pi(s_{t+1})) - Q(s_t, a_t))$$

$$\pi(s) \leftarrow \arg \max_a Q(s, a)$$

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left( r_t + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t) \right)$$

Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$   
Repeat (for each episode):

    Initialize  $S$

    Repeat (for each step of episode):

        Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

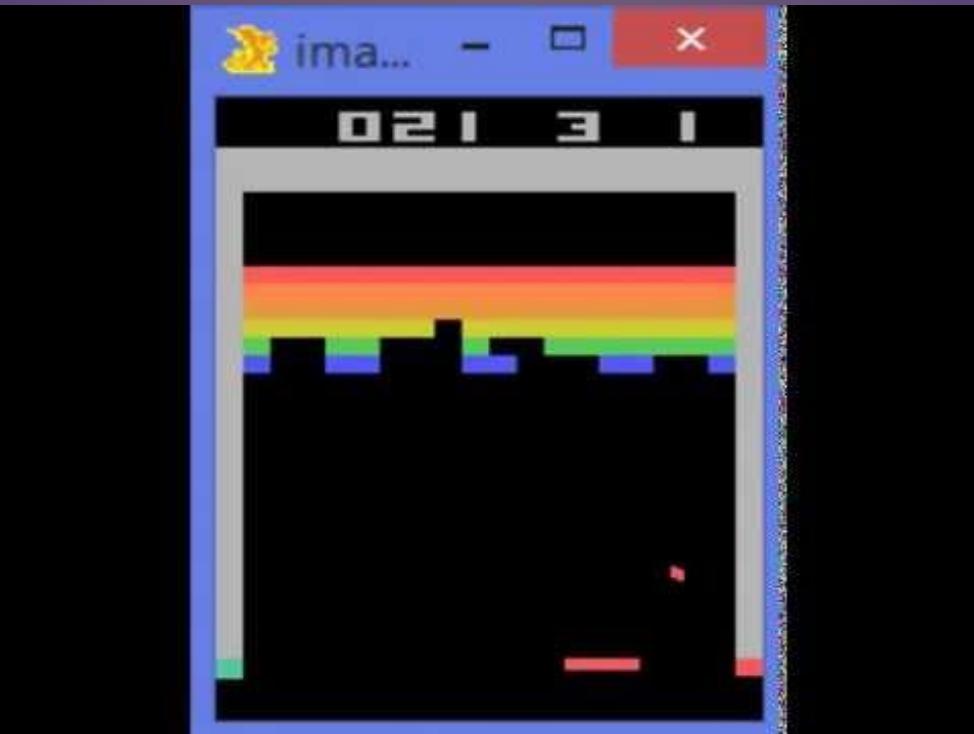
        Take action  $A$ , observe  $R, S'$

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$ ;

    until  $S$  is terminal

# Q-LEARNING





## FUNCTION APPROXIMATION

Model:  $Q_\theta(s_t, a_t)$

Training  
data:

$$\langle s_t, a_t, r_t, s_{t+1} \rangle$$

Loss  
function:

$$\mathcal{L}(\theta) = \|y_t - Q_\theta(s_t, a_t)\|_2^2$$

where  $y_t = r_t + \gamma Q(s_{t+1}, \pi(s_{t+1}))$

# IMPLEMENTATION

## Action-in

$s_t \quad a_t$



$$Q(s_t, a_t)$$

## Action-out

$s_t$



$$Q(s_t, a^{(1)}) \ Q(s_t, a^{(2)}) \ Q(s_t, a^{(i)})$$

## Off-Policy Learning

- The target depends in part on our model → old observations are still useful
- Use a Replay Buffer of most recent transitions as dataset



# Properties of Reinforcement Learning

- **Reward delay**
  - In space invader, only “fire” obtains reward
    - Although the moving before “fire” is important
  - In Go playing, it may be better to sacrifice immediate reward to gain more long-term reward
- Agent’s actions **affect the subsequent data it receives**
  - E.g. Exploration



# DQN ISSUES

→ Convergence is not guaranteed – hope for deep magic!



Replay Buffer

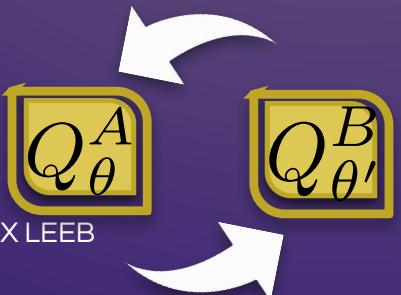


Reward scaling



Using replicas

→ Double Q Learning – decouple action selection and value estimation



$$\theta_B \leftarrow \tau\theta_A + (1 - \tau)\theta_B$$

# POLICY GRADIENTS

- Parameterize policy and update those parameters directly
- Enables new kinds of policies: stochastic, continuous action spaces

$$\cancel{Q_\theta(s, a)} \quad \pi(a|s) \rightarrow \pi_\theta(a|s)$$

- On policy learning → learn directly from your actions



$$\hat{\theta} = \arg \max_{\theta} J(\theta)$$

$$J(\theta) = \mathbb{E}_{a_t \sim \pi(\cdot|s_t;\theta), s_0 \sim \mathcal{T}_0} [G(s_0)]$$

## POLICY GRADIENTS

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \nabla_{\theta} \mathbb{E}_{a_t \sim \pi(\cdot|s_t;\theta)} [G(s_t, a_t)] \\&= \nabla_{\theta} \int \pi(a|s_t;\theta) G(s_t, a) da = \int da G(s_t, a) \nabla_{\theta} \pi(a|s_t;\theta) \\&= \int da G(s_t, a) \frac{\pi(a|s_t;\theta)}{\pi(a|s_t;\theta)} \nabla_{\theta} \pi(a|s_t;\theta) = \int da \pi(a|s_t;\theta) G(s_t, a) \frac{\nabla_{\theta} \pi(a|s_t;\theta)}{\pi(a|s_t;\theta)} \\&= \int da \pi(a|s_t;\theta) G(s_t, a) \nabla_{\theta} \ln \pi(a|s_t;\theta) \\&= \mathbb{E}_{a_t \sim \pi(\cdot|s_t;\theta)} [G(s_t, a_t) \nabla_{\theta} \ln \pi(a_t|s_t;\theta)] \\&\rightarrow \text{Approximate expectation value from samples}\end{aligned}$$

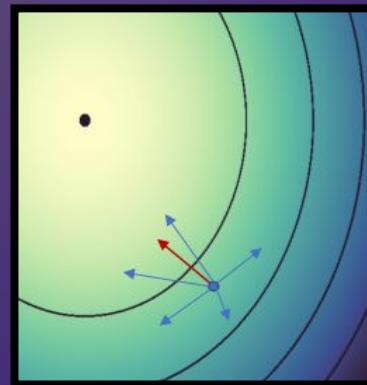
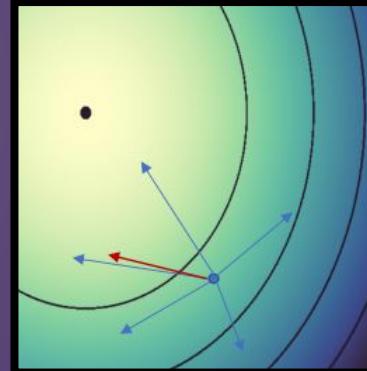
## VARIANCE REDUCTION

- Constant offsets make it harder to differentiate the right direction
- Remove offset → a priori value of each state

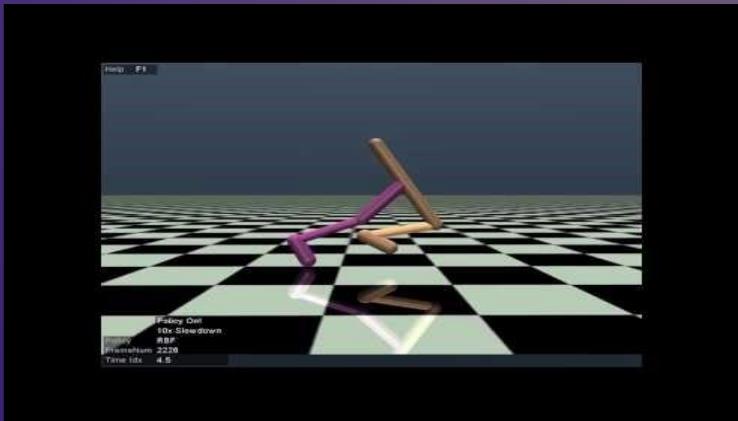
$$G(s_t, a_t) \approx Q(s_t, a_t)$$

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{a_t \sim \pi(\cdot | s_t; \theta)} [(Q(s_t, a_t) - V(s_t)) \nabla_{\theta} \ln \pi(a_t | s_t; \theta)]$$

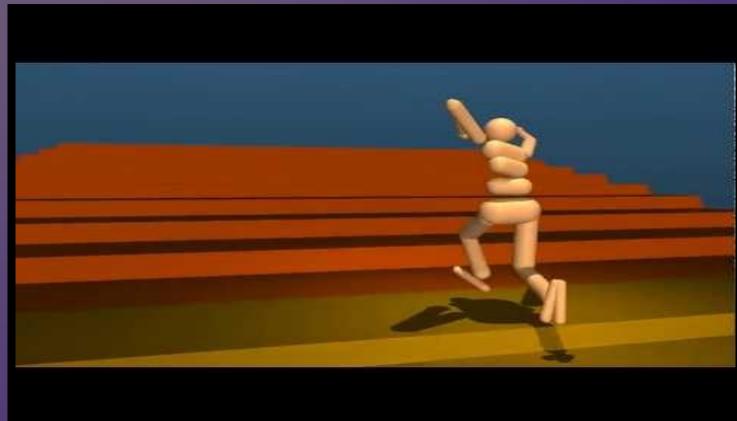
$$\nabla_{\theta} J(\theta) = \mathbb{E}_{a_t \sim \pi(\cdot | s_t; \theta)} [A(s_t, a_t) \nabla_{\theta} \ln \pi(a_t | s_t; \theta)]$$



# ADVANCED POLICY GRADIENT METHODS



Rajeswaran et al.  
(2017)



Heess et al.  
(2017)

ACTOR CRITIC



Critic

using Q learning update

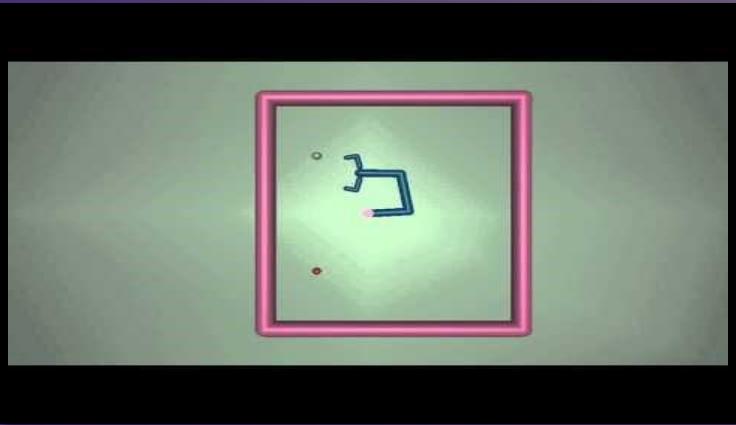
Estimate  
Advantage

Propose  
Actions

Actor

using policy gradient update

# ASYNC ADVANTAGE ACTOR-CRITIC (A3C)



Mnih et al.  
(2016)



# ASYNC ADVANTAGE ACTOR-CRITIC (A3C)





# Deep Reinforcement Learning

## Actor-Critic

# Actor-Critic

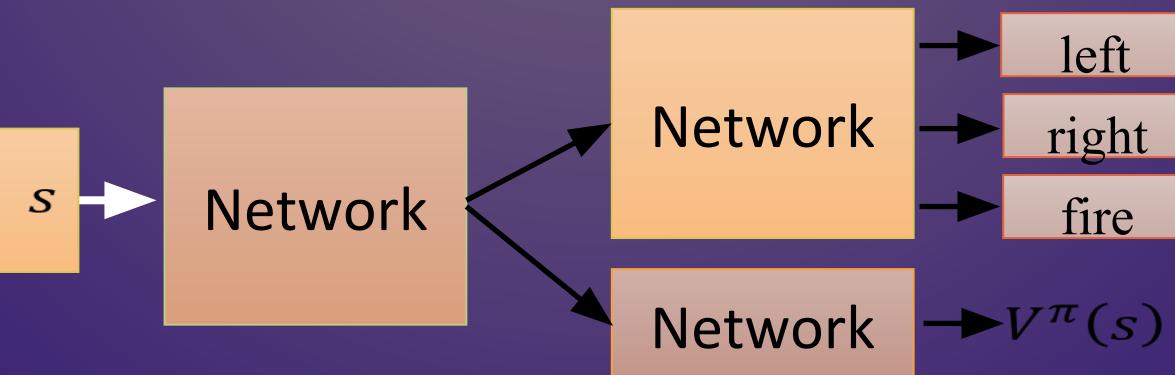
interacts with the environment

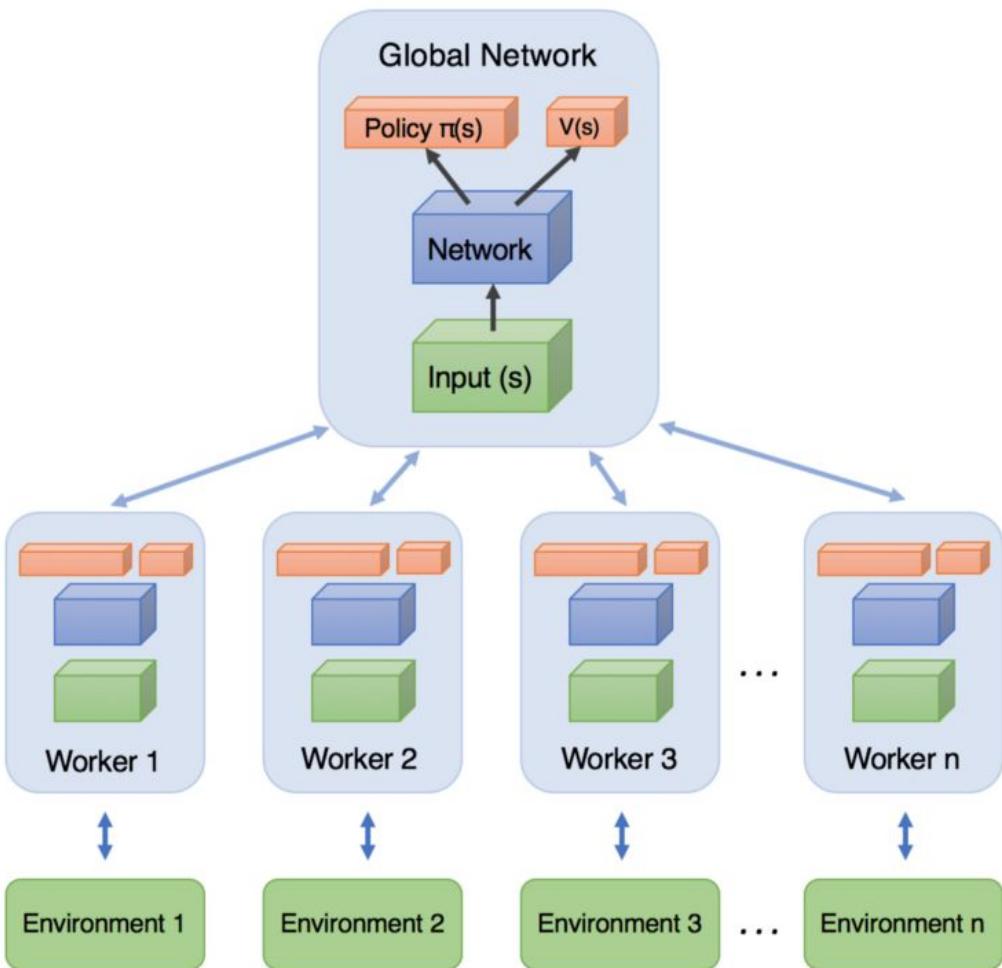
Update actor from critic based on

Learning

# Actor-Critic

- Tips
  - The parameters of actor  $\pi(s)$  and critic  $V^\pi(s)$  can be shared





# Demo of A3C

- Visual Doom AI Competition @ CIG 2016
- <https://www.youtube.com/watch?v=94EPSjQH38Y>

# Why is it challenging

- Exploration-exploitation dilemma
- How to reward the algorithm.
- How to learn when rewards are very sparse
- What representation do we need for states?
- How to update the policy
- How to incorporate the prior (or logic-based) knowledge
- How to learn for multiple tasks: **General Artificial Intelligence**

# Reference

- Textbook: Reinforcement Learning: An Introduction
  - <http://incompleteideas.net/sutton/book/the-book.html>
- Lectures of David Silver
  - <http://www0.cs.ucl.ac.uk/staff/D.Silver/web/Teaching.html> (10 lectures, around 1:30 each)
  - [http://videolectures.net/rldm2015\\_silver\\_reinforcement\\_learning/](http://videolectures.net/rldm2015_silver_reinforcement_learning/) (Deep Reinforcement Learning )
- Lectures of John Schulman
  - [https://youtu.be/aUrX-rP\\_ss4](https://youtu.be/aUrX-rP_ss4)